

# CSCE 633 HW3

March 23, 2020

Submitted by Rizu Jain (430000753)

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
[2]: df_train_raw = pd.read_csv("Train_data.csv")
df_val_raw = pd.read_csv("Validation_data.csv")
df_test_raw = pd.read_csv("Test_data.csv")
```

```
[3]: train_images = df_train_raw['pixels'].str.split(' ', expand=True)
val_images = df_val_raw['pixels'].str.split(' ', expand=True)
test_images = df_test_raw['pixels'].str.split(' ', expand=True)

train_labels = df_train_raw['emotion']
val_labels = df_val_raw['emotion']
test_labels = df_test_raw['emotion']
```

```
[4]: train_images_raw = train_images.apply(pd.to_numeric)
val_images_raw = val_images.apply(pd.to_numeric)
test_images_raw = test_images.apply(pd.to_numeric)
```

## 1 (a) (1 point) Visualization:

```
[5]: emotion_sample_images = {}
fig = plt.figure(figsize=(48,48))
pic_id = 0

for emotion in train_labels.unique():
    images_ids = train_labels.loc[train_labels == emotion].index
    emotion_sample_images[emotion] = images_ids[:5]

    for idx in images_ids[:5]:
        ax = fig.add_subplot(7, 5, pic_id+1, xticks=[], yticks=[])
        ax.imshow(train_images_raw.iloc[idx,:].values.reshape(48, 48),
        cmap='gray')
```

```
ax.set_title(str(emotion), fontsize=30)
pic_id += 1
```



## 2 (b) (0.5 points) Data exploration:

```
[6]: train_labels.value_counts()
```

```
[6]: 3    7215
      6    4965
      4    4830
      2    4097
      0    3995
      5    3171
      1     436
      Name: emotion, dtype: int64
```

```
[7]: from keras.regularizers import l2
      from keras.models import Sequential
      from keras.callbacks import EarlyStopping
      from keras.layers.advanced_activations import PReLU
      from keras.wrappers.scikit_learn import KerasClassifier
      from keras.layers.core import Dense, Dropout, Activation
      from keras.layers.normalization import BatchNormalization
      from sklearn.metrics import accuracy_score
      from sklearn.model_selection import RandomizedSearchCV
      # from sklearn.grid_search import GridSearchCV
      from keras.utils import to_categorical
```

Using TensorFlow backend.

```
[8]: from hyperopt import hp, fmin, tpe, STATUS_OK, Trials
      from keras.layers import Conv2D, Flatten, MaxPooling2D
      from keras import backend as K
      from hyperopt import space_eval

      import time
      import random
```

```
[9]: # pre-processing
      train_images_norm = (train_images_raw / 255) - 0.5
      val_images_norm = (val_images_raw / 255) - 0.5
      test_images_norm = (test_images_raw / 255) - 0.5
```

```
[10]: x_train_fnn = train_images_norm.copy()
      x_val_fnn = val_images_norm.copy()
      x_test_fnn = test_images_norm.copy()
```

```
[11]: y_train = to_categorical(train_labels)
      y_val = to_categorical(val_labels)
      y_test = to_categorical(test_labels)
```

```
[12]: n_input = train_images.shape[1]
      n_class = train_labels.unique().shape[0]
      print('feature num: ', n_input)
```

```
print('class num: ', n_class)
```

feature num: 2304

class num: 7

### 3 (c) (3.5 points) Image classification with FNNs

#### 3.1 (c.i) (2.5 points)

```
[13]: K.clear_session()
```

```
[14]: def optimize_fnn(hyperparameter):
    K.clear_session()
    fnn_model = Sequential()

    first_layer = True
    for layer_size in hyperparameter['network_config']:

        if first_layer:
            fnn_model.add(Dense(layer_size, input_dim = n_input,
→activation=hyperparameter['activation_function'], kernel_regularizer =
→l2(hyperparameter['l2_penalty']),))
            first_layer = False
        else:
            fnn_model.add(Dense(layer_size,
→activation=hyperparameter['activation_function'], kernel_regularizer =
→l2(hyperparameter['l2_penalty'])))

    fnn_model.add(Dropout(hyperparameter['dropout_prob']))

    fnn_model.add(Dense(n_class, activation='softmax'))

    fnn_model.compile(optimizer=hyperparameter['optimizer'],
→loss='categorical_crossentropy', metrics=['accuracy'],)

    train_X, train_y = train_images, y_train
    valid_X, valid_y = val_images, y_val

    start = time.time()
    fnn_history = fnn_model.fit(x_train_fnn, y_train, epochs=20, batch_size=256,
→verbose=0)
    end = time.time()

    # print(fnn_model.summary())

    performance_fnn = fnn_model.evaluate(x_val_fnn, y_val, verbose=0)
```

```

print("-----")

print("Hyperparameters: ", (hyperparameter))
print()
print("No. of parameters : %d" % (fnn_model.count_params()))
print("Training time : %.2f s" % (end-start))
print("Training Accuracy : %.3f" % (fnn_history.history['accuracy'][-1]))
print("Validation Accuracy : %.3f" % (performance_fnn[1]))
print("-----")

if random.random() < 0.35:
    plt.plot(fnn_history.history['loss'])
    plt.title('Cross-entropy loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.show()
    print("-----")

# We want to minimize loss i.e. negative of accuracy
return({"status": STATUS_OK, "loss": -1*performance_fnn[1], "model":fnn_model})

```

```

[15]: # Define search space for hyper-parameters
space_fnn = {
    # The kernel_size for convolutions:
    'network_config': hp.choice('network_config',
→ [[32,32,32], [32,64,128], [64,64], [64,128,64]]),
    # The stride_size for convolutions:
    'activation_function': hp.choice('activation_function', ['tanh', 'relu']),
    # Uniform distribution in finding appropriate dropout values
    'dropout_prob': hp.uniform('dropout_prob', 0.1, 0.35),
    # regularization:
    'l2_penalty': hp.choice('l2_penalty', [0.001,0.003,0.01,0.03]),
    # Choice of optimizer
    'optimizer': hp.choice('optimizer', ['Adam','sgd']),
}

trials_fnn = Trials()

# Find the best hyperparameters
best_fnn = fmin(
    optimize_fnn,
    space_fnn,
    algo=tpe.suggest,
    trials=trials_fnn,
    max_evals=10,
)

```

-----  
Hyperparameters:

```
{'activation_function': 'relu', 'dropout_prob': 0.2658661419495981,  
'l2_penalty': 0.001, 'network_config': (32, 32, 32), 'optimizer': 'Adam'}
```

No. of parameters : 76103

Training time : 38.27 s

Training Accuracy : 0.390

Validation Accuracy : 0.400  
-----

Hyperparameters:

```
{'activation_function': 'relu', 'dropout_prob': 0.34525558508915966,  
'l2_penalty': 0.001, 'network_config': (64, 128, 64), 'optimizer': 'Adam'}
```


No. of parameters : 164551

Training time : 39.13 s

Training Accuracy : 0.406

Validation Accuracy : 0.408  
-----


10%| | 1/10 [01:18<05:54,  
39.36s/trial, best loss: -0.3998328149318695]



-----  
Hyperparameters:

```
{'activation_function': 'tanh', 'dropout_prob': 0.1781299110474086,  
'l2_penalty': 0.03, 'network_config': (32, 32, 32), 'optimizer': 'Adam'}  
No. of parameters : 76103  
Training time : 39.20 s  
Training Accuracy : 0.318  
Validation Accuracy : 0.339
```

-----  
20%| | 2/10 [01:58<05:15,  
39.41s/trial, best loss: -0.407634437084198]



-----  
-----  
Hyperparameters:  
{'activation\_function': 'tanh', 'dropout\_prob': 0.33285994826713994,  
'l2\_penalty': 0.003, 'network\_config': (32, 64, 128), 'optimizer': 'Adam'}  
No. of parameters : 85095  
Training time : 39.04 s  
Training Accuracy : 0.359  
Validation Accuracy : 0.372

-----  
-----  
Hyperparameters:  
{'activation\_function': 'relu', 'dropout\_prob': 0.23281380586559214,  
'l2\_penalty': 0.03, 'network\_config': (32, 64, 128), 'optimizer': 'sgd'}


No. of parameters : 85095  
Training time : 38.39 s  
Training Accuracy : 0.262  
Validation Accuracy : 0.264

-----  
Hyperparameters:

{'activation\_function': 'relu', 'dropout\_prob': 0.2224703044554185,  
'l2\_penalty': 0.003, 'network\_config': (32, 32, 32), 'optimizer': 'Adam'}

No. of parameters : 76103  
Training time : 38.38 s  
Training Accuracy : 0.378  
Validation Accuracy : 0.390

-----  
50%| | 5/10 [03:55<03:15,  
39.20s/trial, best loss: -0.407634437084198]



-----  
Hyperparameters:

{'activation\_function': 'tanh', 'dropout\_prob': 0.33050577093499756,  
'l2\_penalty': 0.03, 'network\_config': (32, 32, 32), 'optimizer': 'Adam'}

No. of parameters : 76103  
Training time : 37.62 s



Training Accuracy : 0.311  
Validation Accuracy : 0.318

-----  
Hyperparameters:

{'activation\_function': 'relu', 'dropout\_prob': 0.22750854237118034,  
'l2\_penalty': 0.03, 'network\_config': (64, 128, 64), 'optimizer': 'sgd'}

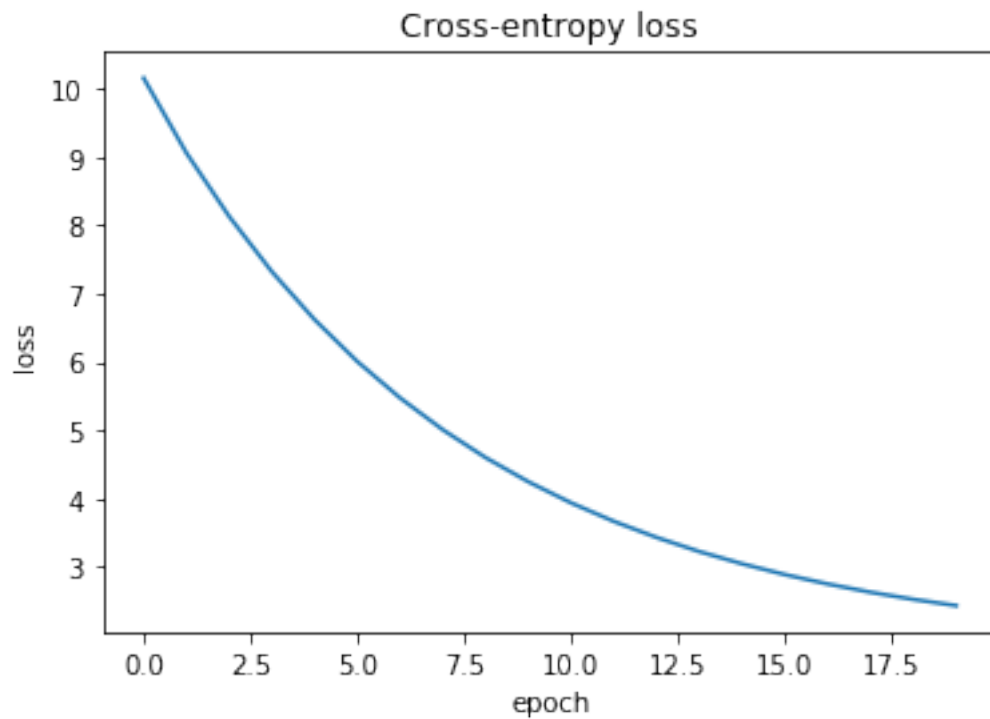
No. of parameters : 164551

Training time : 39.63 s

Training Accuracy : 0.276

Validation Accuracy : 0.284

-----  
70%| | 7/10 [05:13<01:56,  
38.73s/trial, best loss: -0.407634437084198]



-----  
Hyperparameters:

{'activation\_function': 'relu', 'dropout\_prob': 0.17663707686718594,  
'l2\_penalty': 0.03, 'network\_config': (64, 64), 'optimizer': 'sgd'}

No. of parameters : 152135

Training time : 38.40 s

Training Accuracy : 0.335

Validation Accuracy : 0.359

```

-----
Hyperparameters:
{'activation_function': 'relu', 'dropout_prob': 0.2788555859969819,
 'l2_penalty': 0.001, 'network_config': (32, 32, 32), 'optimizer': 'Adam'}
No. of parameters : 76103
Training time : 38.00 s
Training Accuracy : 0.387
Validation Accuracy : 0.405
-----
100%|| 10/10 [06:30<00:00,
39.03s/trial, best loss: -0.407634437084198]

```

### 3.2 (c.ii) (1 point)

```

[16]: print("=====")
      print("Best Hyperparameters", best_fnn)

      best_hyper_fnn = space_eval(space_fnn,best_fnn)
      fnn_tuned = optimize_fnn(best_hyper_fnn)

      performance_fnn = fnn_tuned['model'].evaluate(x_test_fnn, y_test,verbose=0)

      print("=====")
      print("Test Accuracy: ", performance_fnn[1])

```

```

=====
Best Hyperparameters {'activation_function': 1, 'dropout_prob':
0.34525558508915966, 'l2_penalty': 0, 'network_config': 3, 'optimizer': 0}
-----
Hyperparameters: {'activation_function': 'relu', 'dropout_prob':
0.34525558508915966, 'l2_penalty': 0.001, 'network_config': (64, 128, 64),
'optimizer': 'Adam'}

No. of parameters : 164551
Training time : 40.88 s
Training Accuracy : 0.405
Validation Accuracy : 0.408
-----
=====
Test Accuracy: 0.40930622816085815

```

**FNN accuracy on test set: 0.4093**

```

[17]: x_train_cnn = train_images_norm.values.reshape((-1,48,48,1))
      x_val_cnn = val_images_norm.values.reshape((-1,48,48,1))
      x_test_cnn = test_images_norm.values.reshape((-1,48,48,1))

```

## 4 (d) (3 points) Image classification with CNNs:

### 4.1 (d.i) (2 points)

```
[18]: def optimize_cnn(hyperparameter):
    K.clear_session()
    cnn_model = Sequential()

    for layer_sequence in hyperparameter['conv2d_config']:
        for layer_size in layer_sequence:
            cnn_model.add(Conv2D(layer_size,
→kernel_size=hyperparameter['conv_kernel_size'],padding='same',
→strides=hyperparameter['stride_size'],
→activation=hyperparameter['activation_function']))
            cnn_model.add(MaxPooling2D(pool_size=(2,2),padding='same'))
            cnn_model.add(Dropout(hyperparameter['dropout_prob']))

    cnn_model.add(Flatten())
    cnn_model.add(Dense(hyperparameter['dense_layer_neurons'],
→activation=hyperparameter['activation_function']))

    cnn_model.add(Dense(7, activation='softmax'))

    cnn_model.compile(optimizer=hyperparameter['optimizer'],
→loss='categorical_crossentropy', metrics=['accuracy'],)

    start = time.time()
    cnn_history = cnn_model.fit(x_train_cnn, y_train, epochs=20, batch_size=256,
→verbose=0)
    end = time.time()

    # print(cnn_model.summary())

    performance = cnn_model.evaluate(x_val_cnn, y_val, verbose=0)

    print("-----")

    print("Hyperparameters: ", (hyperparameter))
    print()
    print("No. of parameters : %d" % (cnn_model.count_params()))
    print("Training time : %.2f s" % (end-start))
    print("Training Accuracy : %.3f" % (cnn_history.history['accuracy'][-1]))
    print("Validation Accuracy : %.3f" % (performance[1]))
    print("-----")

    if random.random() < 0.35:
        plt.plot(cnn_history.history['loss'])
```

```

plt.title('Cross-entropy loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.show()
print("-----")

# We want to minimize loss i.e. negative of accuracy
return({"status": STATUS_OK, "loss": -1*performance[1], "model":cnn_model})

```

```

[19]: # Define search space for hyper-parameters
space_cnn = {
    # The convolution layers and sizes
    'conv2d_config': hp.choice('conv2d_config',
→ [[(32,32), (64,64)], [(32), (64), (128)]]),
    # The last dense layer size
    'dense_layer_neurons': hp.choice('dense_layer_neurons', [128,256]),
    # The kernel_size for convolutions:
    'conv_kernel_size': hp.choice('conv_kernel_size', [1, 3, 5]),
    # The stride_size for convolutions:
    'stride_size': hp.choice('stride_size', [1, 2]),
    # The activation function
    'activation_function': hp.choice('activation_function', ['tanh', 'relu']),
    # Uniform distribution in finding appropriate dropout values
    'dropout_prob': hp.uniform('dropout_prob', 0.1, 0.35),
    # Choice of optimizer
    'optimizer': hp.choice('optimizer', ['Adam', 'sgd']),
}

trials_cnn = Trials()

# Find the best hyperparameters
best_cnn = fmin(
    optimize_cnn,
    space_cnn,
    algo=tpe.suggest,
    trials=trials_cnn,
    max_evals=10,
)

```


```

-----
Hyperparameters:
{'activation_function': 'tanh', 'conv2d_config': ((32, 32), (64, 64)),
 'conv_kernel_size': 1, 'dense_layer_neurons': 128, 'dropout_prob':
0.14596884626860737, 'optimizer': 'sgd', 'stride_size': 1}
No. of parameters : 1188071
Training time : 848.58 s
Training Accuracy : 0.401

```

Validation Accuracy : 0.378

-----  
0%|  
| 0/10 [14:10<?, ?trial/s, best loss=?]



-----  
Hyperparameters:

{'activation\_function': 'relu', 'conv2d\_config': ((32, 32), (64, 64)),  
'conv\_kernel\_size': 3, 'dense\_layer\_neurons': 256, 'dropout\_prob':  
0.2911961806773443, 'optimizer': 'Adam', 'stride\_size': 1}

No. of parameters : 2426343

Training time : 3183.59 s

Training Accuracy : 0.947

Validation Accuracy : 0.587

-----  
Hyperparameters:

{'activation\_function': 'relu', 'conv2d\_config': ((32, 32), (64, 64)),  
'conv\_kernel\_size': 1, 'dense\_layer\_neurons': 128, 'dropout\_prob':  
0.12362773461891252, 'optimizer': 'Adam', 'stride\_size': 2}

No. of parameters : 16615

Training time : 227.52 s

Training Accuracy : 0.258

Validation Accuracy : 0.270

-----  
Hyperparameters:

{'activation\_function': 'tanh', 'conv2d\_config': ((32,), (64,), (128,)),  
'conv\_kernel\_size': 3, 'dense\_layer\_neurons': 256, 'dropout\_prob':  
0.273851560737502, 'optimizer': 'sgd', 'stride\_size': 1}


No. of parameters : 1274375

Training time : 1466.33 s

Training Accuracy : 0.371

Validation Accuracy : 0.367

-----  
30%| | 3/10 [1:35:38<2:14:46,  
1155.24s/trial, best loss: -0.5865143537521362]



-----  
Hyperparameters:

{'activation\_function': 'tanh', 'conv2d\_config': ((32, 32), (64, 64)),  
'conv\_kernel\_size': 3, 'dense\_layer\_neurons': 256, 'dropout\_prob':  
0.34119556919902994, 'optimizer': 'sgd', 'stride\_size': 2}


No. of parameters : 83431

Training time : 183.41 s

Training Accuracy : 0.290

Validation Accuracy : 0.314

-----  
40%| | 4/10 [1:38:42<2:04:56,  
1249.34s/trial, best loss: -0.5865143537521362]



-----  
Hyperparameters:

{'activation\_function': 'tanh', 'conv2d\_config': ((32, 32), (64, 64)),  
'conv\_kernel\_size': 3, 'dense\_layer\_neurons': 256, 'dropout\_prob':  
0.2818391037926066, 'optimizer': 'sgd', 'stride\_size': 1}


No. of parameters : 2426343

Training time : 2536.79 s

Training Accuracy : 0.417

Validation Accuracy : 0.411

-----  
50%| | 5/10 [2:21:03<1:17:29,  
929.81s/trial, best loss: -0.5865143537521362]



-----  
 -----  
 Hyperparameters:

```
{'activation_function': 'tanh', 'conv2d_config': ((32,), (64,), (128,)),
'conv_kernel_size': 5, 'dense_layer_neurons': 256, 'dropout_prob':
0.13736009970345767, 'optimizer': 'Adam', 'stride_size': 2}
```

No. of parameters : 291847


Training time : 310.60 s

Training Accuracy : 0.815

Validation Accuracy : 0.539

-----  
 60%| | 6/10 [2:26:16<1:34:12,  
 1413.23s/trial, best loss: -0.5865143537521362]





-----  
 -----  
 Hyperparameters:

```
{'activation_function': 'relu', 'conv2d_config': ((32,), (64,), (128,)),
'conv_kernel_size': 5, 'dense_layer_neurons': 128, 'dropout_prob':
0.25213564993880666, 'optimizer': 'Adam', 'stride_size': 2}
```

No. of parameters : 274439

Training time : 362.60 s

Training Accuracy : 0.596

Validation Accuracy : 0.549

-----  
 -----

Hyperparameters:

```
{'activation_function': 'relu', 'conv2d_config': ((32,), (64,), (128,)),
'conv_kernel_size': 3, 'dense_layer_neurons': 256, 'dropout_prob':
0.2008206210895777, 'optimizer': 'sgd', 'stride_size': 2}
```

No. of parameters : 127495

Training time : 294.81 s

Training Accuracy : 0.253

Validation Accuracy : 0.253


-----  
 -----

Hyperparameters:

```
{'activation_function': 'tanh', 'conv2d_config': ((32,), (64,), (128,)),
```

```
'conv_kernel_size': 5, 'dense_layer_neurons': 128, 'dropout_prob':  
0.17873674765950004, 'optimizer': 'Adam', 'stride_size': 2}  
No. of parameters : 274439  
Training time : 367.63 s  
Training Accuracy : 0.760  
Validation Accuracy : 0.551
```

```
-----  
90%|      | 9/10 [2:43:26<11:36,  
696.10s/trial, best loss: -0.5865143537521362]
```



```
-----  
100%|| 10/10 [2:43:27<00:00,  
980.72s/trial, best loss: -0.5865143537521362]
```

```
[20]: print("=====  
print("Best Hyperparameters", best_cnn)  
  
best_hyper_cnn = space_eval(space_cnn,best_cnn)  
cnn_tuned = optimize_cnn(best_hyper_cnn)  
  
performance_cnn = cnn_tuned['model'].evaluate(x_test_cnn, y_test,verbose=0)  
  
print("=====  
print("Test Accuracy: ", performance_cnn[1])
```

=====

Best Hyperparameters {'activation\_function': 1, 'conv2d\_config': 0,  
'conv\_kernel\_size': 1, 'dense\_layer\_neurons': 1, 'dropout\_prob':  
0.2911961806773443, 'optimizer': 0, 'stride\_size': 0}

-----

Hyperparameters: {'activation\_function': 'relu', 'conv2d\_config': ((32, 32),  
(64, 64)), 'conv\_kernel\_size': 3, 'dense\_layer\_neurons': 256, 'dropout\_prob':  
0.2911961806773443, 'optimizer': 'Adam', 'stride\_size': 1}


No. of parameters : 2426343

Training time : 5303.13 s

Training Accuracy : 0.949

Validation Accuracy : 0.582

-----



-----

=====

Test Accuracy: 0.5828921794891357

**CNN accuracy on test set: 0.5829**

## 4.2 FNN and CNN comparision

	FNN	CNN
No. of parameters	< 200K	> 1M
Training time	< 1 min	> 5 min
Best Training accuracy	0.405	0.949
Best Validation accuracy	0.408	0.582
Test accuracy	0.409	0.582

**Conclusion:** We get improvement in our predictions using CNN over FNN.

## 5 (d) (2 points) Fine-tuning:

```
[21]: from keras.applications.vgg16 import VGG16
      from keras.preprocessing import image
      from keras.applications.vgg16 import preprocess_input
      from keras.layers import Input, Flatten, Dense
      from keras.models import Model, load_model
      from keras.datasets import mnist
```

```
[22]: # Convert the images into 3 channels
      x_train_rgb=np.dstack([train_images_norm] * 3)
      x_val_rgb=np.dstack([val_images_norm]*3)
      x_test_rgb=np.dstack([test_images_norm]*3)
      x_train_vgg = x_train_rgb.reshape(-1,48,48,3)
      x_val_vgg= x_val_rgb.reshape (-1,48,48,3)
      x_test_vgg= x_test_rgb.reshape (-1,48,48,3)
```

```
[23]: fer_shape = x_train_vgg.shape[1:]
```

```
[27]: def optimize_vgg(hyperparameter):
      K.clear_session()
      vgg_pretrain = VGG16(weights='imagenet', include_top=False)

      for layer in vgg_pretrain.layers[:hyperparameter['unfreeze_layers']]:
          layer.trainable = False

      fer_input = Input(shape=fer_shape, name = 'image_input')
      vgg_output = vgg_pretrain(fer_input)
      x = Flatten(name='flatten')(vgg_output)

      for layer_size in hyperparameter['add_config']:
          x = Dense(layer_size,
          →activation=hyperparameter['activation_function'],kernel_regularizer =
          →l2(hyperparameter['l2_penalty']))(x)
          x = Dropout(hyperparameter['dropout_prob'])(x)
```

```

x = Dense(n_class, activation='softmax', name='predictions')(x)

fer_pretrained = Model(inputs=fer_input, outputs=x)
fer_pretrained.compile(loss='categorical_crossentropy', optimizer='adam',
↳metrics=['accuracy'])

start = time.time()
vgg_history = fer_pretrained.fit(x_train_vgg, y_train, epochs=2,
↳batch_size=256, verbose=0)
end = time.time()

performance = fer_pretrained.evaluate(x_val_vgg, y_val, verbose=0)

print("-----")

print("Hyperparameters: ", (hyperparameter))
print()
print("No. of parameters : %d" % (fer_pretrained.count_params()))
print("Training time : %.2f s" % (end-start))
print("Training Accuracy : %.3f" % (vgg_history.history['accuracy'][-1]))
print("Validation Accuracy : %.3f" % (performance[1]))
print("-----")

if random.random() < 0.35:
    plt.plot(vgg_history.history['loss'])
    plt.title('Cross-entropy loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.show()
    print("-----")

# We want to minimize loss i.e. negative of accuracy
return({"status": STATUS_OK, "loss": -1*performance[1], "model":
↳fer_pretrained})

```

```

[28]: # Define search space for hyper-parameters
space_vgg = {
    # The convolution layers and sizes
    'add_config': hp.choice('add_config', [[128,32],[256,64]]),
    # The activation function
    'activation_function': hp.choice('activation_function', ['tanh', 'relu']),
    # Uniform distribution in finding appropriate dropout values
    'dropout_prob': hp.uniform('dropout_prob', 0.1, 0.35),
    'unfreeze_layers': hp.choice('unfreeze_layers', [4,5]),
    # regularization:
    'l2_penalty': hp.choice('l2_penalty', [0.001,0.003,0.01,0.03]),
}

```

```

trials_vgg = Trials()

# Find the best hyperparameters
best_vgg = fmin(
    optimize_vgg,
    space_vgg,
    algo=tpe.suggest,
    trials=trials_vgg,
    max_evals=5,
)

```

```

-----
Hyperparameters:
{'activation_function': 'relu', 'add_config': (128, 32), 'dropout_prob':
0.18526718592673108, 'l2_penalty': 0.01, 'unfreeze_layers': 4}
No. of parameters : 14784711
Training time : 530.47 s
Training Accuracy : 0.221
Validation Accuracy : 0.249

```

```

-----
Hyperparameters:
{'activation_function': 'relu', 'add_config': (128, 32), 'dropout_prob':
0.34545365334991956, 'l2_penalty': 0.001, 'unfreeze_layers': 5}
No. of parameters : 14784711
Training time : 512.90 s
Training Accuracy : 0.226
Validation Accuracy : 0.249

```

```

-----
Hyperparameters:
{'activation_function': 'tanh', 'add_config': (128, 32), 'dropout_prob':
0.17201805991803998, 'l2_penalty': 0.003, 'unfreeze_layers': 5}
No. of parameters : 14784711
Training time : 1793.18 s
Training Accuracy : 0.200
Validation Accuracy : 0.249

```

```

-----
Hyperparameters:
{'activation_function': 'relu', 'add_config': (256, 64), 'dropout_prob':
0.24408521902718688, 'l2_penalty': 0.001, 'unfreeze_layers': 4}
No. of parameters : 14862919
Training time : 384.16 s
Training Accuracy : 0.228
Validation Accuracy : 0.249

```

```

-----
Hyperparameters:
{'activation_function': 'tanh', 'add_config': (256, 64), 'dropout_prob':
0.3287684730760788, 'l2_penalty': 0.001, 'unfreeze_layers': 4}
No. of parameters : 14862919
Training time : 378.46 s
Training Accuracy : 0.184
Validation Accuracy : 0.169
-----

```

```

100%|| 5/5 [1:03:22<00:00,
760.42s/trial, best loss: -0.24937307834625244]

```

```

[29]: print("=====")
      print("Best Hyperparameters", best_vgg)

      best_hyper_vgg = space_eval(space_vgg,best_vgg)
      vgg_tuned = optimize_vgg(best_hyper_vgg)

      performance_vgg = vgg_tuned['model'].evaluate(x_test_vgg, y_test,verbose=0)

      print("=====")
      print("Test Accuracy: ", performance_vgg[1])

```

```

=====
Best Hyperparameters {'activation_function': 1, 'add_config': 0, 'dropout_prob':
0.18526718592673108, 'l2_penalty': 2, 'unfreeze_layers': 0}
-----

```

```

Hyperparameters: {'activation_function': 'relu', 'add_config': (128, 32),
'dropout_prob': 0.18526718592673108, 'l2_penalty': 0.01, 'unfreeze_layers': 4}

No. of parameters : 14784711
Training time : 404.27 s
Training Accuracy : 0.242
Validation Accuracy : 0.249
-----

```

```

=====
Test Accuracy: 0.24491502344608307

```

**Fine-Tuning accuracy on the test set: 0.2449**