

CSCE 633 Homework 4

April 9, 2020

Submitted by Rizu Jain (430000753)

```
[1]: import math

import pandas as pd
import numpy as np

from sklearn import tree
from sklearn.metrics import mean_squared_error
```

```
[2]: # Load Dataset

train_data = pd.read_csv('OnlineNewsPopularityTrain.csv')
test_data = pd.read_csv('OnlineNewsPopularityTest.csv')

# Remove label column
train_arr = train_data.values[:,1:]
test_arr = test_data.values[:,1:]
```

1 Decision tree regression

```
[3]: # Prepare randomly shuffled data for decision tree

np.random.shuffle(train_arr)
np.random.shuffle(test_arr)

x_train = train_arr[:, :-1]
y_train = train_arr[:, -1]
x_test = test_arr[:, :-1]
y_test = test_arr[:, -1]
```

```
[4]: samples_per_fold= x_train.shape[0]//5

kfold_x_train=[]
kfold_y_train=[]
for i in range(5):
    kfold_x_train.append(x_train[samples_per_fold*i:samples_per_fold*(i+1),:])
```

```
kfold_y_train.append(y_train[samples_per_fold*i:samples_per_fold*(i+1)])
```

```
[5]: depths= [2,5,3,7,10,20]
depth_err = []

for curr_depth in depths:
    crossval_err = []
    for i in range(5):
        clf = tree.DecisionTreeRegressor(max_depth=curr_depth,random_state=0)

        train_xsets = []
        train_ysets = []
        curr_x_test = []
        curr_y_test = []

        for j in range(5):
            if j==i:
                curr_x_test = kfold_x_train[i]
                curr_y_test = kfold_y_train[i]
            else:
                train_xsets.append(kfold_x_train[j])
                train_ysets.append(kfold_y_train[j])

        curr_x_train = np.
→concatenate((train_xsets[0],train_xsets[1],train_xsets[2],train_xsets[3]))
        curr_y_train = np.
→concatenate((train_ysets[0],train_ysets[1],train_ysets[2],train_ysets[3]))

        clf = clf.fit(curr_x_train,curr_y_train)
        curr_y_pred = clf.predict(curr_x_test)

        err = mean_squared_error(curr_y_pred,curr_y_test)
        crossval_err.append(math.sqrt(err))

    print("Cross validation errors for max depth : ",curr_depth)
    print(crossval_err)
    depth_err.append(np.mean(crossval_err))
    print()

print("-----")
for k in range(len(depths)):
    print("Max depth: ",depths[k],"\tCrossVal average RSS: ",depth_err[k])
```

```
Cross validation errors for max depth : 2
[9224.11604620015, 12065.381162979264, 8510.082656879917, 15842.621659452852,
```

11418.842315740498]

Cross validation errors for max depth : 5

[14306.989870526284, 12179.479621233402, 8405.625688241142, 17406.20798278216, 11360.913452065786]

Cross validation errors for max depth : 3

[14030.46728927809, 12196.151465152392, 8717.023898819332, 17403.776402822554, 11463.558051403981]

Cross validation errors for max depth : 7

[14929.8187040767, 14386.506099851586, 15001.675912877789, 19062.406462796498, 11606.225755964482]

Cross validation errors for max depth : 10

[15841.729347189994, 13388.82413275757, 17386.360791621697, 20290.145438720225, 15758.206878103683]

Cross validation errors for max depth : 20

[16978.09970117997, 16327.387296521709, 17816.150254216573, 21102.489465270446, 13923.153016805458]

```
-----  
Max depth: 2   CrossVal average RSS: 11412.208768250535  
Max depth: 5   CrossVal average RSS: 12731.843322969755  
Max depth: 3   CrossVal average RSS: 12762.19542149527  
Max depth: 7   CrossVal average RSS: 14997.32658711341  
Max depth: 10  CrossVal average RSS: 16533.053317678634  
Max depth: 20  CrossVal average RSS: 17229.45594679883
```

[6]: *# Performance on Test data*

```
optimal_depth=depths[depth_err.index(min(depth_err))]  
clf = tree.DecisionTreeRegressor(max_depth=optimal_depth,random_state=0)  
clf = clf.fit(x_train,y_train)  
y_pred = clf.predict(x_test)  
  
test_err = mean_squared_error(y_pred,y_test)  
test_err = math.sqrt(test_err)  
  
print("Optimal Depth: ",optimal_depth)  
print("Test Error: ",test_err)
```

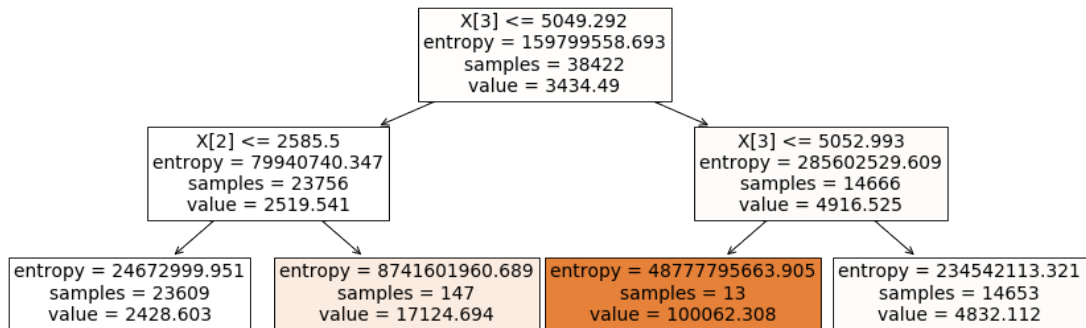
Optimal Depth: 2

Test Error: 8296.439356745595

2 Feature Exploration

```
[12]: import matplotlib.pyplot as plt

plt.rcParams['figure.figsize'] = [15, 5]
tree.plot_tree(clf,filled = True);
```



```
[8]: # we dumped the 1st feature from our original dataframe
imp_features = [27,30,28]

print("The most important features are :")
for feature in imp_features:
    print(train_data.columns[feature])
```

The most important features are :

- kw_avg_avg
- self_reference_avg_sharess
- self_reference_min_shares

27. kw_avg_avg: Avg. keyword (avg. shares)
28. self_reference_min_shares: Min. shares of referenced articles in Mashable
30. self_reference_avg_sharess: Avg. shares of referenced articles in Mashable

Intuition

1. The use of keywords is an important factor in predicting number of shares
2. The popularity of referenced articles is another important factor

3 Random Forest

```
[9]: # hyperparamters to tune
# No. of trees
# depth of trees

num_trees_options = [10,20,40,80]
tree_depth = [2,4,8,16]
tree_depth_err = []

for num_trees in num_trees_options:
    depth_err = []
    for curr_depth in tree_depth:

        crossval_err = []

        for i in range(5):
            train_xsets = []
            train_ysets = []
            curr_x_test = []
            curr_y_test = []

            for j in range(5):
                if j==i:
                    curr_x_test = kfold_x_train[i]
                    curr_y_test = kfold_y_train[i]
                else:
                    train_xsets.append(kfold_x_train[j])
                    train_ysets.append(kfold_y_train[j])

            curr_x_train = np.
→concatenate((train_xsets[0],train_xsets[1],train_xsets[2],train_xsets[3]))
            curr_y_train = np.
→concatenate((train_ysets[0],train_ysets[1],train_ysets[2],train_ysets[3]))

            predictions= np.zeros(curr_y_test.shape)

            for _ in range(num_trees):
                # select data samples randomly
                # we select as many samples as original but we allow repeats
                # source: https://stats.stackexchange.com/questions/347818/
→number-of-samples-per-tree-in-a-random-forest

                sample_ids = np.random.choice(curr_x_train.shape[0],curr_x_train.
→shape[0])
```

```

        x_train_rf = curr_x_train[sample_ids]
        y_train_rf = curr_y_train[sample_ids]

        # randomly select 30% features for this tree
        sel_columns = np.random.choice(x_train_rf.
→shape[1],int(x_train_rf.shape[1]*0.3))
        x_train_rf = x_train_rf[:,sel_columns]

        clf = tree.DecisionTreeRegressor(max_depth=curr_depth)
        clf = clf.fit(x_train_rf,y_train_rf)
        curr_y_pred = clf.predict(curr_x_test[:,sel_columns])
        predictions += curr_y_pred

    predictions = predictions/num_trees
    err = mean_squared_error(predictions,curr_y_test)
    crossval_err.append(math.sqrt(err))

    print("-----")
    print("max depth : ",curr_depth)
    print("num trees : ",num_trees)
    print("crossval_err : ",crossval_err)
    depth_err.append(np.mean(crossval_err))
    print()

tree_depth_err.append(depth_err)

```

```

-----
max depth : 2
num trees : 10
crossval_err : [9185.340878314053, 12070.276586605261, 8412.881754990376,
15799.307737160054, 11387.441569869328]

```

```

-----
max depth : 4
num trees : 10
crossval_err : [9437.13326781748, 12136.258396977764, 8586.780269075025,
15791.31208518285, 11379.1349516635]

```

```

-----
max depth : 8
num trees : 10
crossval_err : [9883.822881673846, 12305.610623246203, 8987.271770430261,
16259.022486803888, 11750.999839408036]

```

```

-----
max depth : 16

```

num trees : 10
crossval_err : [9970.775065148686, 12753.854540515482, 9190.986467418059,
16120.390661007328, 12304.681234486488]

max depth : 2
num trees : 20
crossval_err : [9198.277123395585, 12085.858670854968, 8372.829357713827,
15748.975697556178, 11351.79005153414]

max depth : 4
num trees : 20
crossval_err : [9224.842218391972, 12141.02336328227, 8380.66491990371,
15773.471087204663, 11380.044039262653]

max depth : 8
num trees : 20
crossval_err : [9410.835143283497, 12171.924615451335, 8547.666218286065,
15766.101157368608, 11504.103902963549]

max depth : 16
num trees : 20
crossval_err : [9530.990573779967, 12337.639951767642, 8821.508866663818,
15981.598200205788, 11948.794700090133]

max depth : 2
num trees : 40
crossval_err : [9211.763895639724, 12068.790062841606, 8370.885648698306,
15752.059079752073, 11349.154047842516]

max depth : 4
num trees : 40
crossval_err : [9251.393064615833, 12098.39339174754, 8399.667006008976,
15783.759791873154, 11346.256497292554]

max depth : 8
num trees : 40
crossval_err : [9377.762785446423, 12117.381635336376, 8501.161040856203,
15807.077836383245, 11442.9907904012]

max depth : 16

```
num trees : 40
crossval_err : [9489.313675408115, 12244.529619787018, 8618.082989713932,
15842.955591464131, 11513.829050615977]
```

```
-----
max depth : 2
num trees : 80
crossval_err : [9199.377005245928, 12075.617294139556, 8377.0216163679,
15781.51826091811, 11343.19267337088]
```

```
-----
max depth : 4
num trees : 80
crossval_err : [9216.044340993618, 12082.352916696347, 8375.696474494987,
15752.374612246558, 11347.832993223]
```

```
-----
max depth : 8
num trees : 80
crossval_err : [9339.208114168956, 12094.450458670784, 8419.930332885371,
15759.189305968373, 11380.786447905291]
```

```
-----
max depth : 16
num trees : 80
crossval_err : [9273.454408219484, 12112.916673038406, 8495.880566165053,
15810.977699166513, 11444.682154373224]
```

```
[14]: # 2-dimensional color-coded matrix
# x/y dimensions are the number of trees and tree depth, and
# the color-coding reflects the average error over all folds


import seaborn as sns

xlabels = [str(i) for i in tree_depth ]
ylabels = [str(j) for j in num_trees_options ]

ax = sns.heatmap(tree_depth_err,
                  annot=True,
                  fmt = "0.2f",
                  cmap= "coolwarm",
                  annot_kws={'size':16},
                  xticklabels=xlabels,
                  yticklabels=ylabels,
                  robust=True,
                  linewidths=.5)
```



```
ax.set(xlabel='Tree Depth', ylabel='Number of Trees',title="Random Forest Cross_
validation RMSE")
plt.show()
```



```
[11]: # Performance on Test Set

err_random_forest = np.asarray(tree_depth_err)
ind = np.unravel_index(np.argmin(err_random_forest, axis=None),
    →err_random_forest.shape)
best_num_trees = num_trees_options[ind[0]]
best_tree_depth = tree_depth[ind[1]]

final_predictions= np.zeros(y_test.shape)

for _ in range(best_num_trees):
    # select data samples randomly
    # we select as many samples as original but we allow repeats
    # source: https://stats.stackexchange.com/questions/347818/
    →number-of-samples-per-tree-in-a-random-forest

    sample_ids = np.random.choice(x_train.shape[0],x_train.shape[0])
    features_rf = x_train[sample_ids]
    labels_rf = y_train[sample_ids]

    # randomly select 30% features for this tree
    sel_columns = np.random.choice(features_rf.shape[1],int(features_rf.
    →shape[1]*0.3))
    features_rf = features_rf[:,sel_columns]

    clf = tree.DecisionTreeRegressor(max_depth=best_tree_depth)
```

```

    clf = clf.fit(features_rf, labels_rf)
    y_pred = clf.predict(x_test[:, sel_columns])
    final_predictions += y_pred

final_predictions = final_predictions / best_num_trees
test_err = mean_squared_error(final_predictions, y_test)
test_err = math.sqrt(test_err)

print("-----")
print("Max depth : ", curr_depth)
print("Bum trees : ", num_trees)
print("Test Error : ", test_err)
print("-----")

```

```

-----
Max depth :  16
Bum trees :  80
Test Error : 8337.195445959338
-----

```

Test Error Random Forest : 8337.2